

DDA5001 Machine Learning

Lec 2: Basics and Review

Junchi Yang

School of Data Science
The Chinese University of Hong Kong, Shenzhen



Basic Mathematics

Concepts of Learning

Basic Notions of Linear Algebra

Basic Notions of Linear Algebra

- **Vector**. $\mathbf{x} \in \mathbb{R}^n$ is a real-valued n -dimensional **column** vector; i.e.,

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad x_i \in \mathbb{R} \ \forall i.$$

- You can regard the vector $\mathbf{x} \in \mathbb{R}^n$ as a point in the n -dimensional **linear space** $\mathbb{R}^n$ (Think of $n = 2$ and $n = 3$).

Basic Notions of Linear Algebra

- **Vector**. $\mathbf{x} \in \mathbb{R}^n$ is a real-valued n -dimensional **column** vector; i.e.,

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad x_i \in \mathbb{R} \ \forall i.$$

- You can regard the vector $\mathbf{x} \in \mathbb{R}^n$ as a point in the n -dimensional **linear space** $\mathbb{R}^n$ (Think of $n = 2$ and $n = 3$).
- **Addition of vectors**. The addition of two vectors is defined by adding corresponding coordinates, i.e.,

$$\begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} + \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} x_1 + y_1 \\ \vdots \\ x_n + y_n \end{bmatrix},$$

Basic Notions of Linear Algebra

- **Multiplication.** The multiplication of a scalar with a vector is defined by performing multiplication in each coordinate:

$$a \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} ax_1 \\ \vdots \\ ax_n \end{bmatrix},$$

where $a \in \mathbb{R}$.

Basic Notions of Linear Algebra

- **Multiplication.** The multiplication of a scalar with a vector is defined by performing multiplication in each coordinate:

$$a \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} ax_1 \\ \vdots \\ ax_n \end{bmatrix},$$

where $a \in \mathbb{R}$.

- **Commutativity.** $\mathbf{x} + \mathbf{y} = \mathbf{y} + \mathbf{x}$ for all $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$.
- **Distributive properties** $a(\mathbf{x} + \mathbf{y}) = a\mathbf{x} + a\mathbf{y}$ and $(a + b)\mathbf{x} = a\mathbf{x} + b\mathbf{x}$ for all $a, b \in \mathbb{R}$ and $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$.
- **Transpose of vector.** Let

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^n.$$

The transpose of $\mathbf{x}$ is

$$\mathbf{x}^\top = [x_1 \quad x_2 \quad \cdots \quad x_n].$$

Basic Notions of Linear Algebra

- **Linear independence.** We say that a finite collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ of vectors in $\mathbb{R}^n$ is **linearly dependent** if there exist scalars $a_1, \dots, a_m \in \mathbb{R}$, **not all of them are zero**, such that

$$\sum_{i=1}^m a_i \mathbf{x}_i = \mathbf{0}.$$

The collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is said to be **linearly independent** if it is **not** linearly dependent.

Basic Notions of Linear Algebra

- **Linear independence.** We say that a finite collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ of vectors in $\mathbb{R}^n$ is **linearly dependent** if there exist scalars $a_1, \dots, a_m \in \mathbb{R}$, **not all of them are zero**, such that

$$\sum_{i=1}^m a_i \mathbf{x}_i = \mathbf{0}.$$

The collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is said to be **linearly independent** if it is **not** linearly dependent.

- **Span.** The set of all linear combinations of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is called the span of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$, i.e.,
 $\text{span}\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\} := \left\{ \sum_{i=1}^m a_i \mathbf{x}_i : \mathbf{a} \in \mathbb{R}^m \right\}$

Basic Notions of Linear Algebra

- **Linear independence.** We say that a finite collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ of vectors in $\mathbb{R}^n$ is **linearly dependent** if there exist scalars $a_1, \dots, a_m \in \mathbb{R}$, **not all of them are zero**, such that

$$\sum_{i=1}^m a_i \mathbf{x}_i = \mathbf{0}.$$

The collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is said to be **linearly independent** if it is **not** linearly dependent.

- **Span.** The set of all linear combinations of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is called the span of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$, i.e.,
 $\text{span}\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\} := \left\{ \sum_{i=1}^m a_i \mathbf{x}_i : \mathbf{a} \in \mathbb{R}^m \right\}$
- **Basis.** A basis of a vector space $\mathcal{V}$ is a collection of vectors in $\mathcal{V}$ that is linearly independent and spans $\mathcal{V}$.

Basic Notions of Linear Algebra

- **Linear independence.** We say that a finite collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ of vectors in $\mathbb{R}^n$ is **linearly dependent** if there exist scalars $a_1, \dots, a_m \in \mathbb{R}$, **not all of them are zero**, such that

$$\sum_{i=1}^m a_i \mathbf{x}_i = \mathbf{0}.$$

The collection $C = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is said to be **linearly independent** if it is **not** linearly dependent.

- **Span.** The set of all linear combinations of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$ is called the span of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\}$, i.e.,
 $\text{span}\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m\} := \left\{ \sum_{i=1}^m a_i \mathbf{x}_i : \mathbf{a} \in \mathbb{R}^m \right\}$
- **Basis.** A basis of a vector space $\mathcal{V}$ is a collection of vectors in $\mathcal{V}$ that is linearly independent and spans $\mathcal{V}$. For example,

$$\left\{ \begin{bmatrix} 1 \\ 2 \end{bmatrix}, \begin{bmatrix} 3 \\ 4 \end{bmatrix} \right\} \quad \text{and} \quad \left\{ \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \end{bmatrix} \right\}$$

are bases of $\mathbb{R}^2$.

Basic Notions of Linear Algebra

- **Inner product.** Given two vectors $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{y} \in \mathbb{R}^n$, their inner product is defined as

$$\langle \mathbf{x}, \mathbf{y} \rangle = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i$$

We say that $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ are **orthogonal** if $\mathbf{x}^\top \mathbf{y} = 0$.

Basic Notions of Linear Algebra

- **Inner product.** Given two vectors $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{y} \in \mathbb{R}^n$, their inner product is defined as

$$\langle \mathbf{x}, \mathbf{y} \rangle = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i$$

We say that $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ are **orthogonal** if $\mathbf{x}^\top \mathbf{y} = 0$.

- **Euclidean/ ℓ_2 -norm.** For vector $\mathbf{x} = [x_1 \ x_2 \ \cdots \ x_n]^\top \in \mathbb{R}^n$,

$$\|\mathbf{x}\|_2 = \sqrt{\mathbf{x}^\top \mathbf{x}} = \sqrt{\sum_{i=1}^n x_i^2},$$

which measures the length of $\mathbf{x}$. For simplicity, we often only write $\|\mathbf{x}\|$ to represent $\|\mathbf{x}\|_2$.

Basic Notions of Linear Algebra

- **Inner product.** Given two vectors $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{y} \in \mathbb{R}^n$, their inner product is defined as

$$\langle \mathbf{x}, \mathbf{y} \rangle = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i$$

We say that $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ are **orthogonal** if $\mathbf{x}^\top \mathbf{y} = 0$.

- **Euclidean/ ℓ_2 -norm.** For vector $\mathbf{x} = [x_1 \ x_2 \ \cdots \ x_n]^\top \in \mathbb{R}^n$,

$$\|\mathbf{x}\|_2 = \sqrt{\mathbf{x}^\top \mathbf{x}} = \sqrt{\sum_{i=1}^n x_i^2},$$

which measures the length of $\mathbf{x}$. For simplicity, we often only write $\|\mathbf{x}\|$ to represent $\|\mathbf{x}\|_2$.

- More generally, a **norm** $\|\cdot\| : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function that satisfies
- $\|\mathbf{x}\| > 0$ for all $\mathbf{x} \neq 0$ and $\|\mathbf{x}\| = 0$ only if $\mathbf{x} = 0$;
 - $\|a\mathbf{x}\| = |a|\|\mathbf{x}\|$ for $\mathbf{x} \in \mathbb{R}^n$ and $a \in \mathbb{R}$;
 - $\|\mathbf{x} + \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\|$ for all $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ (**triangle inequality**)

Basic Notions of Linear Algebra

- **Hölder p -norm.** We now introduce common norms in $\mathbb{R}^n$ —the Hölder p -norm, $1 \leq p \leq \infty$, which are defined as follows:

$$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

for $1 \leq p < \infty$ and

$$\|\mathbf{x}\|_\infty = \max_{1 \leq i \leq n} |x_i|.$$

Basic Notions of Linear Algebra

- **Hölder p -norm.** We now introduce common norms in $\mathbb{R}^n$ —the Hölder p -norm, $1 \leq p \leq \infty$, which are defined as follows:

$$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

for $1 \leq p < \infty$ and

$$\|\mathbf{x}\|_\infty = \max_{1 \leq i \leq n} |x_i|.$$

- **Special cases.** When $p = 2$, it reduces to the ℓ_2 -norm. When $p = 1$, it reduces to the ℓ_1 -norm, i.e.,

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|.$$

Basic Notions of Linear Algebra

- **Hölder p -norm.** We now introduce common norms in $\mathbb{R}^n$ —the Hölder p -norm, $1 \leq p \leq \infty$, which are defined as follows:

$$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

for $1 \leq p < \infty$ and

$$\|\mathbf{x}\|_\infty = \max_{1 \leq i \leq n} |x_i|.$$

- **Special cases.** When $p = 2$, it reduces to the ℓ_2 -norm. When $p = 1$, it reduces to the ℓ_1 -norm, i.e.,

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|.$$

- **Cauchy-Schwarz inequality.**

$$|\mathbf{x}^\top \mathbf{y}| \leq \|\mathbf{x}\|_2 \|\mathbf{y}\|_2 \quad \forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^n.$$

Basic Notions of Linear Algebra

- **Matrix.** We use $\mathbb{R}^{m \times n}$ to denote the set of $m \times n$ arrays whose components are from $\mathbb{R}$. We can write a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ as

$$\mathbf{A} = \begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{m1} & \cdots & a_{mn} \end{bmatrix}, \quad a_{i,j} \in \mathbb{R} \quad \forall i, j.$$

Basic Notions of Linear Algebra

- **Matrix.** We use $\mathbb{R}^{m \times n}$ to denote the set of $m \times n$ arrays whose components are from $\mathbb{R}$. We can write a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ as

$$\mathbf{A} = \begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{m1} & \cdots & a_{mn} \end{bmatrix}, \quad a_{i,j} \in \mathbb{R} \quad \forall i, j.$$

- **Transpose of Matrix.** Given an $m \times n$ matrix $\mathbf{A}$, its transpose $\mathbf{A}^\top$ is defined as the following $n \times m$ matrix:

$$\mathbf{A}^\top = \begin{bmatrix} a_{11} & a_{21} & \cdots & a_{m1} \\ a_{12} & a_{22} & \cdots & a_{m2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \cdots & a_{mn} \end{bmatrix},$$

Basic Notions of Linear Algebra

- **Matrix.** We use $\mathbb{R}^{m \times n}$ to denote the set of $m \times n$ arrays whose components are from $\mathbb{R}$. We can write a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ as

$$\mathbf{A} = \begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{m1} & \cdots & a_{mn} \end{bmatrix}, \quad a_{i,j} \in \mathbb{R} \quad \forall i, j.$$

- **Transpose of Matrix.** Given an $m \times n$ matrix $\mathbf{A}$, its transpose $\mathbf{A}^\top$ is defined as the following $n \times m$ matrix:

$$\mathbf{A}^\top = \begin{bmatrix} a_{11} & a_{21} & \cdots & a_{m1} \\ a_{12} & a_{22} & \cdots & a_{m2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \cdots & a_{mn} \end{bmatrix},$$

- **Symmetric matrix.** An $m \times m$ real matrix $\mathbf{A}$ is said to be symmetric if $\mathbf{A} = \mathbf{A}^\top$.

Basic Notions of Linear Algebra

- **Matrix-matrix multiplication.** The matrix-matrix multiplication between $\mathbf{A} \in \mathbb{R}^{m \times n}$ and $\mathbf{B} \in \mathbb{R}^{n \times p}$ is defined as

$$\mathbb{R}^{m \times p} \ni \mathbf{C} = \mathbf{AB} \quad \text{where} \quad c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}.$$

Illustration:

$$\begin{bmatrix} c_{11} & c_{12} & c_{13} \\ c_{21} & c_{22} & c_{23} \\ c_{31} & c_{32} & c_{33} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \end{bmatrix}$$

The matrix-vector multiplication can be viewed as a special case of matrix-matrix multiplication, i.e., with $\mathbf{A} \in \mathbb{R}^{m \times n}$ and $\mathbf{b} \in \mathbb{R}^n$, we have

$$\mathbb{R}^m \ni \mathbf{c} = \mathbf{Ab} \quad \text{where} \quad c_i = \sum_{k=1}^n a_{ik} b_k.$$

Basic Notions of Linear Algebra

- ▶ **Three perspectives for matrix-matrix multiplication.** There are three (equivalent) important ways for interpreting $C = AB$:

Basic Notions of Linear Algebra

- ▶ **Three perspectives for matrix-matrix multiplication.** There are three (equivalent) important ways for interpreting $C = AB$:

- The first one is by definition

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}, \quad \forall i = 1, 2, \dots, m. \quad j = 1, 2, \dots, p.$$

Basic Notions of Linear Algebra

- **Three perspectives for matrix-matrix multiplication.** There are three (equivalent) important ways for interpreting $C = AB$:

- The first one is by definition

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}, \quad \forall i = 1, 2, \dots, m. \quad j = 1, 2, \dots, p.$$

- The second one is by **outer product**

$$C = \sum_{k=1}^n \mathbf{a}_k \mathbf{b}_k^{\top},$$

where $\mathbf{a}_k$ and $\mathbf{b}_k^{\top}$ are the k -th column and row of A and B , respectively.

Basic Notions of Linear Algebra

- **Three perspectives for matrix-matrix multiplication.** There are three (equivalent) important ways for interpreting $C = AB$:

- The first one is by definition

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}, \quad \forall i = 1, 2, \dots, m. \quad j = 1, 2, \dots, p.$$

- The second one is by **outer product**

$$C = \sum_{k=1}^n \mathbf{a}_k \mathbf{b}_k^{\top},$$

where $\mathbf{a}_k$ and $\mathbf{b}_k^{\top}$ are the k -th column and row of A and B , respectively.

- The third one is by matrix-vector product

$$\mathbf{c}_j = A \mathbf{b}_j, \quad \forall j = 1, 2, \dots, p.$$

Basic Notions of Linear Algebra

- **Rank.** The rank of a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$, denoted by $\text{rank}(\mathbf{A})$, is defined as the number of elements of a maximal linearly independent subset of its columns or rows.

Basic Notions of Linear Algebra

- ▶ **Rank.** The rank of a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$, denoted by $\text{rank}(\mathbf{A})$, is defined as the number of elements of a maximal linearly independent subset of its columns or rows.
- ▶ **Matrix inverse.** An $n \times n$ square matrix $\mathbf{A}$ is said to be invertible if the columns of $\mathbf{A}$ has full-rank. The inverse of the matrix $\mathbf{A}$ is denoted as $\mathbf{A}^{-1}$, and we have

$$\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}.$$

Facts:

- $(\mathbf{A}^{-1})^{-1} = \mathbf{A}$.
- $(\mathbf{AB})^{-1} = \mathbf{B}^{-1}\mathbf{A}^{-1}$, where $\mathbf{A}, \mathbf{B}$ are square and invertible.

Basic Notions of Linear Algebra

- **Orthogonal matrix.** An $n \times n$ square matrix $\mathbf{A}$ is said to be orthogonal, or orthonormal, is a real square matrix whose columns and rows are orthonormal vectors. That is,

$$\mathbf{A}^\top \mathbf{A} = \mathbf{A} \mathbf{A}^\top = \mathbf{I}$$

In another word, for orthogonal matrix $\mathbf{A}$, we have

$$\mathbf{A}^\top = \mathbf{A}^{-1}.$$

Basic Notions of Linear Algebra

- **Orthogonal matrix.** An $n \times n$ square matrix $\mathbf{A}$ is said to be orthogonal, or orthonormal, is a real square matrix whose columns and rows are orthonormal vectors. That is,

$$\mathbf{A}^\top \mathbf{A} = \mathbf{A} \mathbf{A}^\top = \mathbf{I}$$

In another word, for orthogonal matrix $\mathbf{A}$, we have

$$\mathbf{A}^\top = \mathbf{A}^{-1}.$$

- **Positive semi-definite (PSD) matrix.** An $n \times n$ real symmetric matrix $\mathbf{A}$ is said to be positive semi-definite if

$$\mathbf{x}^\top \mathbf{A} \mathbf{x} \geq 0, \quad \forall \mathbf{x} \in \mathbb{R}^n.$$

- **Positive definite (PD) matrix.** An $n \times n$ real symmetric matrix $\mathbf{A}$ is said to be positive definite if

$$\mathbf{x}^\top \mathbf{A} \mathbf{x} > 0, \quad \forall \mathbf{x} \in \mathbb{R}^n \setminus \{0\}.$$

Basic Notions of Multivariate Calculus

Basic Notions of Multivariate Calculus

- **Gradient.** It is a generalization of derivative to multi-dimensional functions. Assume $f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$ is continuously differentiable. Then, we denote the gradient of f by (an $n \times 1$ vector):

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Basic Notions of Multivariate Calculus

- **Gradient.** It is a generalization of derivative to multi-dimensional functions. Assume $f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$ is continuously differentiable. Then, we denote the gradient of f by (an $n \times 1$ vector):

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Facts:

- If $f(\mathbf{x}) = \mathbf{c}^\top \mathbf{x}$, then $\nabla f(\mathbf{x}) = \mathbf{c}$.
- If $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{M} \mathbf{x}$ ($\mathbf{M}$ is symmetric), then: $\nabla f(\mathbf{x}) = 2\mathbf{M} \mathbf{x}$.

Basic Notions of Multivariate Calculus

- **Gradient.** It is a generalization of derivative to multi-dimensional functions. Assume $f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$ is continuously differentiable. Then, we denote the gradient of f by (an $n \times 1$ vector):

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$


Facts:

- If $f(\mathbf{x}) = \mathbf{c}^\top \mathbf{x}$, then $\nabla f(\mathbf{x}) = \mathbf{c}$.
- If $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{M} \mathbf{x}$ ($\mathbf{M}$ is symmetric), then: $\nabla f(\mathbf{x}) = 2\mathbf{M} \mathbf{x}$.
- **First-order Taylor expansion.** The first-order Taylor expansion yields:

$$f(\mathbf{y}) = f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{y} - \mathbf{x}) + o(\|\mathbf{y} - \mathbf{x}\|), \quad \|\mathbf{y} - \mathbf{x}\| \rightarrow 0.$$

Basic Notions of Multivariate Calculus

Illustration of first-order Taylor expansion:



- Accurate when y is close to x ; the approximation can be poor farther away.
- Important notion for later first-order algorithm development.

Basic Notions of Probability and Statistics

Basic Notions of Probability and Statistics

- **Expectation.** Suppose X is a random variable, its expectation is denoted as

$$\mathbb{E}[X].$$

Suppose X takes discrete values $x_1, \dots, x_k$ with probability $p_1, \dots, p_k$, then

$$\mathbb{E}[X] = \sum_{i=1}^k p_i x_i.$$

Suppose X takes continuous values in $(-\infty, +\infty)$ with density $p(x)$, then

$$\mathbb{E}[X] = \int_{-\infty}^{+\infty} p(x)x dx.$$

- **Variance.** Suppose X is a random variable, its variance is denoted as

$$\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2].$$

Basic Notions of Probability and Statistics

- ▶ **Random vector.** $\mathbf{X} = [X_1, \dots, X_n]^\top$ is a random vector if each coordinate is a random variable.
- ▶ **Expectation of random vector.** Suppose $\mathbf{X}$ is an n -dimensional random vector, its expectation is denoted as

$$\mathbb{E}[\mathbf{X}] = [\mathbb{E}[X_1], \dots, \mathbb{E}[X_n]]^\top.$$

- ▶ **Covariance matrix.** Suppose $\mathbf{X} = [X_1, \dots, X_n]^\top$ is an n -dimensional random vector, its covariance matrix is $n \times n$ matrix defined as

$$\text{Var}[\mathbf{X}] = \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^\top].$$

Basic Notions of Probability and Statistics

- **Gaussian distribution.** A random variable X is said to follow $\mathcal{N}(\mu, \sigma^2)$ (Gaussian distribution with mean μ and variance σ^2) if its probability density function (PDF) is given by

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right)$$



Basic Notions of Probability and Statistics

- **Gaussian distribution.** A random variable X is said to follow $\mathcal{N}(\mu, \sigma^2)$ (Gaussian distribution with mean μ and variance σ^2) if its probability density function (PDF) is given by

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right)$$

- **Multivariate Gaussian distribution.** We say the random vector $\mathbf{X} \in \mathbb{R}^d$ follows Gaussian distribution with mean $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$ (assumed to be PD), if its PDF is given by

$$p(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2}} \frac{1}{|\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$



Basic Notions of Optimization

- **Optimization**. The seek of maximum or minimum. Formally speaking, finding the minimum value of f over $\mathbb{R}^n$ is written as

$$\min_{\boldsymbol{\theta} \in \mathbb{R}^n} f(\boldsymbol{\theta}).$$

Basic Notions of Optimization


- **Optimization.** The seek of maximum or minimum. Formally speaking, finding the minimum value of f over $\mathbb{R}^n$ is written as

$$\min_{\theta \in \mathbb{R}^n} f(\theta).$$

- **Global minimizer.** Find the point θ^* (called global minimizer/global optimum/optimal solution) that achieves the minimum value of f over $\mathbb{R}^n$

$$\theta^* = \operatorname{argmin}_{\theta \in \mathbb{R}^n} f(\theta).$$

Clearly, $f(\theta^*) = \min_{\theta \in \mathbb{R}^n} f(\theta)$.



Basic Mathematics

Concepts of Learning

Components of Supervised Learning: Motivation from An Example

Learning Example: Credit Approval

Task: Learning to predict if one applicant should be approved a credit card.

Applicant's info.:

age	25
gender	male
salary	100000 RMB
citizenship	CN
years in job	2 year
⋮	⋮

Question: should we approve credit card to the applicant?

Learning Example: Credit Approval

Task: Learning to predict if one applicant should be approved a credit card.

Applicant's info.:

age	25
gender	male
salary	100000 RMB
citizenship	CN
years in job	2 year
⋮	⋮

Question: should we approve credit card to the applicant?

How to automate such a task by using machine learning methods?

Data: Samples

- Collect a series of **historical data**

	age	gender	salary	citizenship	years in job
Applicant 1	2.5	1	10	3	1
Applicant 2	2.8	0	8	6	5
Applicant 3	1.6	0	0	4	0
Applicant 4	2.3	1	8	2	4
Applicant 5	3	0	4	2	1

We have the **data matrix** (**feature matrix**) $\mathbf{X}$ as

$$\mathbf{X} = \begin{bmatrix} 2.5 & 1 & 10 & 3 & 1 \\ 2.8 & 0 & 8 & 6 & 5 \\ 1.6 & 0 & 0 & 4 & 0 \\ 2.3 & 1 & 8 & 2 & 4 \\ 3 & 0 & 4 & 2 & 1 \end{bmatrix}$$

- Each row $\mathbf{x}_i^\top$ is called a **sample**, representing i -th applicant's data.
- Each column is called a **feature**, representing all the applicant's behavior about the j -th feature.

Data: Labels

- Collect the corresponding label

	age	gender	salary	citizenship	years in job	approve
App. 1	2.5	1	10	3	1	+1
App. 2	2.8	0	8	6	5	+1
App. 3	1.6	0	0	4	0	-1
App. 4	2.3	1	8	2	4	+1
App. 5	3	0	4	2	1	-1

We have the label y as

$$y = \begin{bmatrix} +1 \\ +1 \\ -1 \\ +1 \\ -1 \end{bmatrix}$$

- Each y_i represents the label of the i -th applicant.
- Label implies supervision, a very important quantity in machine learning.

Supervised Learning: Hypothesis/Model

- ▶ We have an underlying and unknown hypothesis/model $g \in \mathcal{H}$

$$g : \mathcal{X} \mapsto \mathcal{Y}$$

where $\mathcal{X}$ is the input space (set of all possible inputs), while $\mathcal{Y}$ is the output space (label space).

In our example, g is the target function that maps $\mathbf{x}_i$ to y_i .

- ▶ Learn a model f from the hypothesis/model space $\mathcal{H}$ based on the training dataset $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$.

Ideally, f should fully capture the patterns in data, i.e., it approximates well the target function g

$$f \approx g.$$

- ▶ The hypothesis space $\mathcal{H}$ is one of the hardest parts to be pre-determined in a learning process. One typical instance of $\mathcal{H}$ is the set of all possible linear fit to the data (results in linear models), while another popular choice is nonlinear model (e.g., neural networks).

Supervised Learning: Hypothesis/Model

Parametrization:

$f = f_{\theta} \in \mathcal{H}$ is often parameterized by the parameters θ

Example:

- ▶ In linear regression, $f_{\theta}(x) = \theta^{\top} x$ is all possible linear fits and θ is the parameters of the model. A specific θ determines a specific model.
- ▶ In deep learning, f_{θ} is the neural network and θ represents weights (network parameters), respectively.

Supervised Learning: Hypothesis/Model

Parametrization:

$f = f_{\theta} \in \mathcal{H}$ is often parameterized by the parameters θ

Example:

- ▶ In linear regression, $f_{\theta}(x) = \theta^{\top} x$ is all possible linear fits and θ is the parameters of the model. A specific θ determines a specific model.
- ▶ In deep learning, f_{θ} is the neural network and θ represents weights (network parameters), respectively.

Two main categories of hypothesis space $\mathcal{H}$:

- ▶ **Linear**
 - ▶ Linear regression
 - ▶ Linear classification
- ▶ **Nonlinear**
 - ▶ Neural networks

Supervised Learning: Learning Problem and Algorithm

- ▶ Given training dataset $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$.
- ▶ Choose the hypothesis f_{θ} .
- ▶ Choose the **loss function** $\ell : \mathbb{R} \rightarrow \mathbb{R}$.
- ▶ **Learning/optimization problem**

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(\mathbf{x}_i), y_i) \quad (\mathbf{P})$$

$\rightsquigarrow$ **Optimization algorithm** $\mathcal{A}$ is designed to solve $(\mathbf{P})$.


$\rightsquigarrow$ After learning to obtain $\hat{\theta}$, we get the **learned model** $f_{\hat{\theta}}$. Then, one can use the learned $f_{\hat{\theta}}$ to do prediction.

Supervised Learning: Components

Formalization:

- ▶ Target function $g : \mathcal{X} \rightarrow \mathcal{Y}$ (underlying credit approval model)
- ▶ Training dataset: $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ (historical records)
- ▶ Hypothesis space $\mathcal{H}$ (learning scope to approximate g)
- ▶ Hypothesis/model: f_{θ} (model to be determined)
- ▶ Optimization algorithm: $\mathcal{A}$ (learning the model from data)

Supervised Learning: High-level View



Hopefully,

$$f_{\hat{\theta}} \approx g$$

Predict/decision: When a new sample data (**test data**) x comes, the label is predicted as:

$$y \leftarrow f_{\hat{\theta}}(x).$$

$\rightsquigarrow$ Next lecture: Linear classification and linear regression.